

E534 - Big Data Applications

Lecture Notes

Geoffrey C. Fox
Gregor von Laszewski

Editor

laszewski@gmail.com

<https://cloudmesh-community.github.io/book/vonlaszewski-e534.epub>

September 27, 2019 - 10:05 AM

Created by Cloudmesh & Cyberaide Bookmanager, <https://github.com/cyberaide/bookmanager>

E534 - BIG DATA APPLICATIONS

Geoffrey C. Fox Gregor von Laszewski

(c) Indiana University, Gregor von Laszewski, Geoffrey Fox, 2018, 2019

E534 - BIG DATA APPLICATIONS

1 PREFACE

1.1 Disclaimer

1.1.1 Acknowledgment

1.1.2 Extensions

2 WEEK 1

2.1 Part I Motivation I

2.1.1 Motivation

2.1.2 00) Mechanics of Course, Summary, and overall remarks on course

2.1.2.1 01A) Technology Hypecycle I

2.1.2.2 01B) Technology Hypecycle II

2.1.2.3 01C) Technology Hypecycle III

2.1.2.4 01D) Technology Hypecycle IV

2.1.3 02)

2.1.3.1 02A) Clouds/Big Data Applications I

2.1.3.2 02B) Cloud/Big Data Applications II

2.1.3.3 02C) Cloud/Big Data

2.1.4 03) Jobs In areas like Data Science, Clouds and Computer Science and Computer

2.1.5 04) Industry, Technology, Consumer Trends Basic trends 2018 Lectures 4A 4B have

2.1.6 05) Digital Disruption and Transformation The Past displaced by Digital

2.1.7 06)

2.1.8 06A) Computing Model I Industry adopted clouds which are attractive for data

2.1.8.1 06B) Computing Model II with 3 subsections is removed; please see 2018

2.1.9 07) Research Model 4th Paradigm; From Theory to Data driven science?

2.1.10 08) Data Science Pipeline DIKW: Data, Information, Knowledge, Wisdom, Decisions.

2.1.11 09) Physics: Looking for Higgs Particle with Large Hadron Collider LHC Physics as a big data example

2.1.12 10) Recommender Systems I General remarks and Netflix

example

2.1.13 11) Recommender Systems II Exploring Data Bags and Spaces

2.1.14 12) Web Search and Information Retrieval Another Big Data Example

2.1.15 13) Cloud Applications in Research Removed Science Clouds, Internet of Things

2.1.16 14) Parallel Computing and MapReduce Software Ecosystems

2.1.17 15) Online education and data science education Removed.

2.1.18 16) Conclusions

3 WEEK 2

3.1 Part II Motivation Archive ☁

3.1.1 2018 BDAA Motivation-1A) Technology Hypecycle I

3.1.2 2018 BDAA Motivation-1B) Technology Hypecycle II

3.1.3 2018 BDAA Motivation-2B) Cloud/Big Data Applications II

3.1.4 2018 BDAA Motivation-4A) Industry Trends I

3.1.5 2018 BDAA Motivation-4B) Industry Trends II

3.1.6 2017 BDAA Motivation-4C) Industry Trends III

3.1.7 2018 BDAA Motivation-6B) Computing Model II

3.1.8 2017 BDAA Motivation-8) Data Science Pipeline DIKW

3.1.9 2017 BDAA Motivation-13) Cloud Applications in Research Science Clouds Internet of Things

3.1.10 2017 BDAA Motivation-15) Data Science Education Opportunities at Universities

4 WEEK 3

4.1 Part III Cloud ☁

4.1.1 A. Summary of Course

4.1.2 B. Defining Clouds I

4.1.3 C. Defining Clouds II

4.1.4 D. Defining Clouds III: Cloud Market Share

4.1.5 E. Virtualization: Virtualization Technologies,

4.1.6 F. Cloud Infrastructure I

4.1.7 G. Cloud Infrastructure II

4.1.8 H. Cloud Software:

4.1.9 I. Cloud Applications I: Clouds in science where area called

4.1.10 J. Cloud Applications II: Characterize Applications using NIST

4.1.11 K. Parallel Computing

4.1.12 L. Real Parallel Computing: Single Program/Instruction Multiple

Data SIMD SPMD

[4.1.13 M. Storage: Cloud data](#)

[4.1.14 N. HPC and Clouds](#)

[4.1.15 O. Comparison of Data Analytics with Simulation:](#)

[4.1.16 P. The Future I](#)

[4.1.17 Q. other Issues II](#)

[4.1.18 R. The Future and other Issues III](#)

5 WEEK 4

[5.1 Google Colab](#) 

[5.1.1 How to create a Google Colab Notebook?](#)

[5.1.2 How to create a simple programme with Colab?](#)

[5.1.3 How to do a benchmark with Cloudmesh Tools using Colab?](#)

6 WEEK 5

[6.1 Week 5](#) 

[6.1.1 Deep Learning Lesson 1: MNIST Classification Version 1](#)

[6.1.2 Using Cloudmesh Common](#)

[6.1.3 Import Libraries](#)

[6.1.4 Pre-process data](#)

[6.1.4.1 Load data](#)

[6.1.4.2 Identify Number of Classes](#)

[6.1.4.3 Convert Labels To One-Hot Vector](#)

[6.1.5 Image Reshaping](#)

[6.1.6 Resize and Normalize](#)

[6.1.7 Create a Keras Model](#)

[6.1.8 Compile and Train](#)

[6.1.9 Testing](#)

[6.1.10 Final Note](#)

[6.1.10.1 Reference:](#)

7 ASSIGNMENTS

[7.1 Assignments](#) 

7.2 WEEKLY ASSIGNMENTS

[7.2.1 Assignment 1](#) 

[7.2.2 Assignment 2](#) 

[7.2.3 Assignment 3](#) 

8 GITHUB

[8.1 Track Progress with Github](#) 

[8.1.1 How to check this?](#)

[8.1.1.1 Step 1](#)

[8.1.1.2 Step 2](#)

[8.1.1.3 Step 3](#)

[8.1.1.4 Step 4](#)

[8.1.1.5 Step 5 \(Optional\)](#)

[8.1.1.6 Step 6 \(Optional\)](#)

[9 REFERENCES](#)

1 PREFACE

Fri Sep 27 10:05:22 EDT 2019 [cloud](#)

1.1 DISCLAIMER [cloud](#)

This book has been generated with [Cyberaide Bookmanager](#).

Bookmanager is a tool to create a publication from a number of sources on the internet. It is especially useful to create customized books, lecture notes, or handouts. Content is best integrated in markdown format as it is very fast to produce the output.

Bookmanager has been developed based on our experience over the last 3 years with a more sophisticated approach. Bookmanager takes the lessons from this approach and distributes a tool that can easily be used by others.

The following shields provide some information about it. Feel free to click on them.

1.1.1 Acknowledgment

If you use bookmanager to produce a document you must include the following acknowledgement.

“This document was produced with Cyberaide Bookmanager developed by Gregor von Laszewski available at <https://pypi.python.org/pypi/cyberaide-bookmanager>. It is in the responsibility of the user to make sure an author acknowledgement section is included in your document. Copyright verification of content included in a book is responsibility of the book editor.”

The bibtex entry is

```
@Misc{www-cyberaide-bookmanager,
```

```
author =  {Gregor von Laszewski},  
title =   {{Cyberaide Book Manager}},  
howpublished = {pypi},  
month =    apr,  
year =     2019,  
url={https://pypi.org/project/cyberaide-bookmanager/}  
}
```

1.1.2 Extensions

We are happy to discuss with you bugs, issues and ideas for enhancements.
Please use the convenient github issues at

- <https://github.com/cyberaide/bookmanager/issues>

Please do not file with us issues that relate to an editors book. They will provide you with their own mechanism on how to correct their content.

2 WEEK 1


2.1 PART I MOTIVATION I

2.1.1 Motivation

Big Data Applications & Analytics: Motivation/Overview; Machine (actually Deep) Learning, Big Data, and the Cloud; Centerpieces of the Current and Future Economy,

2.1.2 00) Mechanics of Course, Summary, and overall remarks on course


In this section we discuss the summary of the motivation section.



[Video](#)

2.1.2.1 01A) Technology Hypecycle I

Today clouds and big data have got through the hype cycle (they have emerged) but features like blockchain, serverless and machine learning are on recent hype cycles while areas like deep learning have several entries (as in fact do clouds) Gartner's Hypecycles and especially that for emerging technologies in 2019 The phases of hypecycles Priority Matrix with benefits and adoption time Initial discussion of 2019 Hypecycle for Emerging Technologies



[Video](#)

2.1.2.2 01B) Technology Hypecycle II

Today clouds and big data have got through the hype cycle (they have emerged)

but features like blockchain, serverless and machine learning are on recent hype cycles while areas like deep learning have several entries (as in fact do clouds) Gartner's Hypecycles and especially that for emerging technologies in 2019 Details of 2019 Emerging Technology and related (AI, Cloud) Hypecycles



[Video](#)

2.1.2.3 01C) Technology Hypecycle III


Today clouds and big data have got through the hype cycle (they have emerged) but features like blockchain, serverless and machine learning are on recent hype cycles while areas like deep learning have several entries (as in fact do clouds) Gartners Hypecycles and Priority Matrices for emerging technologies in 2018, 2017 and 2016 More details on 2018 will be found in Unit 1A of 2018 Presentation and details of 2015 in Unit 1B (Journey to Digital Business). 1A in 2018 also discusses 2017 Data Center Infrastructure removed as this hype cycle disappeared in later years.



[Video](#)

2.1.2.4 01D) Technology Hypecycle IV

Today clouds and big data have got through the hype cycle (they have emerged) but features like blockchain, serverless and machine learning are on recent hype cycles while areas like deep learning have several entries (as in fact do clouds) Emerging Technologies hypecycles and Priority matrix at selected times 2008-2015 Clouds star from 2008 to today They are mixed up with transformational and disruptive changes Unit 1B of 2018 Presentation has more details of this history including Priority matrices




[Video](#)

2.1.3 02)

2.1.3.1 02A) Clouds/Big Data Applications I


The Data Deluge Big Data; a lot of the best examples have NOT been updated (as I can't find updates) so some slides old but still make the correct points Big Data Deluge has become the Deep Learning Deluge Big Data is an agreed fact; Deep Learning still evolving fast but has stream of successes!



[Video](#)

2.1.3.2 02B) Cloud/Big Data Applications II

Clouds in science where area called cyberinfrastructure; The usage pattern from NIST is removed. See 2018 lectures 2B of the motivation for this discussion



[Video](#)

2.1.3.3 02C) Cloud/Big Data

Usage Trends Google and related Trends Artificial Intelligence from Microsoft, Gartner and Meeker



[Video](#)

2.1.4 03) Jobs In areas like Data Science, Clouds and Computer Science and Computer

Engineering



[Video](#)

2.1.5 04) Industry, Technology, Consumer Trends Basic trends 2018 Lectures 4A 4B have


more details removed as dated but still valid See 2018 Lesson 4C for 3 Technology trends for 2016: Voice as HCI, Cars, Deep Learning



[Video](#)

2.1.6 05) Digital Disruption and Transformation The Past displaced by Digital

Disruption; some more details are in 2018 Presentation Lesson 5



[Video](#)

2.1.7 06)

2.1.8 06A) Computing Model I Industry adopted clouds which are attractive for data

analytics. Clouds are a dominant force in Industry. Examples are given

2.1.8.1 06B) Computing Model II with 3 subsections is removed; please see 2018

Presentation for this Developments after 2014 mainly from Gartner Cloud Market share Blockchain



[Video](#)


2.1.9 07) Research Model 4th Paradigm; From Theory to Data driven science?



[Video](#)


2.1.10 08) Data Science Pipeline DIKW: Data, Information, Knowledge, Wisdom, Decisions.

More details on Data Science Platforms are in 2018 Lesson 8 presentation



[Video](#)

2.1.11 09) Physics: Looking for Higgs Particle with Large Hadron Collider LHC Physics as a big data example



[Video](#)

2.1.12 10) Recommender Systems I General remarks and Netflix example



[Video](#)

2.1.13 11) Recommender Systems II Exploring Data Bags and Spaces



[Video](#)


2.1.14 12) Web Search and Information Retrieval Another Big Data Example



[Video](#)

2.1.15 13) Cloud Applications in Research Removed Science Clouds, Internet of Things

Part 12 continuation. See 2018 Presentation (same as 2017 for lesson 13) and Cloud Unit 2019-I) this year



[Video](#)

2.1.16 14) Parallel Computing and MapReduce Software Ecosystems



[Video](#)

2.1.17 15) Online education and data science education Removed.

You can find it in the 2017 version. In [Section 3.1](#) you can see more about this.



[Video](#)

2.1.18 16) Conclusions

Conclusion contain in the latter part of the part 15.

Motivation Archive Big Data Applications & Analytics: Motivation/Overview; Machine (actually Deep) Learning, Big Data, and the Cloud; Centerpieces of the Current and Future Economy. Backup Lectures from previous years referenced in 2019 class




[Video](#)

3 WEEK 2

3.1 PART II MOTIVATION ARCHIVE

3.1.1 2018 BDAA Motivation-1A) Technology Hypecycle I


In this section we discuss on general remarks including Hype curves.



[Video](#)

3.1.2 2018 BDAA Motivation-1B) Technology Hypecycle II

In this section we continue our discussion on general remarks including Hype curves.



[Video](#)

3.1.3 2018 BDAA Motivation-2B) Cloud/Big Data Applications II

In this section we discuss clouds in science where area called cyberinfrastructure; the usage pattern from NIST Artificial Intelligence from Gartner and Meeker.




[Video](#)

3.1.4 2018 BDAA Motivation-4A) Industry Trends I

In this section we discuss on Lesson 4A many technology trends through end of


2014.



[Video](#)

3.1.5 2018 BDAA Motivation-4B) Industry Trends II


In this section we continue our discussion on industry trends. This section includes Lesson 4B 2015 onwards many technology adoption trends.



[Video](#)

3.1.6 2017 BDAA Motivation-4C) Industry Trends III


In this section we continue our discussion on industry trends. This section contains lesson 4C 2015 onwards 3 technology trends voice as HCI cars deep learning.



[Video](#)

3.1.7 2018 BDAA Motivation-6B) Computing Model II


In this section we discuss computing models. This section contains lesson 6B with 3 subsections developments after 2014 mainly from Gartner cloud market share blockchain



[Video](#)

3.1.8 2017 BDAA Motivation-8) Data Science Pipeline DIKW


In this section, we discuss data science pipelines. This section also contains about data, information, knowledge, wisdom forming DIKW term. And also it contains some discussion on data science platforms.



[Video](#)

3.1.9 2017 BDAA Motivation-13) Cloud Applications in Research Science Clouds Internet of Things


In this section we discuss about internet of things and related cloud applications.



[Video](#)

3.1.10 2017 BDAA Motivation-15) Data Science Education Opportunities at Universities

In this section we discuss more on data science education opportunities.




[Video](#)

4 WEEK 3

4.1 PART III CLOUD


4.1.1 A. Summary of Course



[Video](#)

4.1.2 B. Defining Clouds I

In this lecture we discuss the basic definition of cloud and two very simple examples of why virtualization is important.



[Video](#)

In this lecture we discuss how clouds are situated wrt HPC and supercomputers, why multicore chips are important in a typical data center.

4.1.3 C. Defining Clouds II

In this lecture we discuss service-oriented architectures, Software services as Message-linked computing capabilities.



[Video](#)

In this lecture we discuss different aaS's: Network, Infrastructure, Platform, Software. The amazing services that Amazon AWS and Microsoft Azure have Initial Gartner comments on clouds (they are now the norm) and evolution of

servers; serverless and microservices Gartner hypecycle and priority matrix on Infrastructure Strategies.

4.1.4 D. Defining Clouds III: Cloud Market Share



[Video](#)

In this lecture we discuss on how important the cloud market shares are and how much money do they make.

4.1.5 E. Virtualization: Virtualization Technologies,



[Video](#)

In this lecture we discuss hypervisors and the different approaches KVM, Xen, Docker and Openstack.

4.1.6 F. Cloud Infrastructure I



[Video](#)

In this lecture we comment on trends in the data center and its technologies. Clouds physically spread across the world Green computing Fraction of world's computing ecosystem. In clouds and associated sizes an analysis from Cisco of size of cloud computing is discussed in this lecture.


4.1.7 G. Cloud Infrastructure II



[Video](#)

In this lecture, we discuss Gartner hypecycle and priority matrix on Compute Infrastructure Containers compared to virtual machines The emergence of artificial intelligence as a dominant force.

4.1.8 H. Cloud Software:



[Video](#)

In this lecture we discuss, HPC-ABDS with over 350 software packages and how to use each of 21 layers Google's software innovations MapReduce in pictures Cloud and HPC software stacks compared Components need to support cloud/distributed system programming.


4.1.9 I. Cloud Applications I: Clouds in science where area called



[Video](#)

In this lecture we discuss cyberinfrastructure; the science usage pattern from NIST Artificial Intelligence from Gartner.


4.1.10 J. Cloud Applications II: Characterize Applications using NIST



[Video](#)

In this lecture we discuss the approach Internet of Things with different types of MapReduce.

4.1.11 K. Parallel Computing



[Video](#)

In this lecture we discuss analogies, parallel computing in pictures and some useful analogies and principles.

4.1.12 L. Real Parallel Computing: Single Program/Instruction Multiple Data SIMD SPMD



[Video](#)

In this lecture, we discuss Big Data and Simulations compared and we furthermore discusses what is hard to do.


4.1.13 M. Storage: Cloud data



[Video](#)

In this lecture we discuss about the approaches, repositories, file systems, data lakes.


4.1.14 N. HPC and Clouds



[Video](#)

In this lecture we discuss the Branscomb Pyramid Supercomputers versus clouds Science Computing Environments.

4.1.15 O. Comparison of Data Analytics with Simulation:



[Video](#)

In this lecture we discuss the structure of different applications for simulations and Big Data Software implications Languages.

4.1.16 P. The Future I



[Video](#)

In this lecture we discuss Gartner cloud computing hypecycle and priority matrix 2017 and 2019 Hyperscale computing Serverless and FaaS Cloud Native Microservices Update to 2019 Hypecycle.

4.1.17 Q. other Issues II



[Video](#)

In this lecture we discuss on Security Blockchain.

4.1.18 R. The Future and other Issues III



[Video](#)

In this lecture we discuss on Fault Tolerance.

5 WEEK 4

5.1 GOOGLE COLAB

In this section we are going to introduce you, how to use Google Colab to run deep learning models.

5.1.1 How to create a Google Colab Notebook?



5.1.2 How to create a simple programme with Colab?



5.1.3 How to do a benchmark with Cloudmesh Tools using Colab?



6 WEEK 5

6.1 WEEK 5

6.1.1 Deep Learning Lesson 1: MNIST Classification Version 1

6.1.2 Using Cloudmesh Common

Here we do a simple benchmark. We calculate compile time, train time, test time and data loading time for this example. Installing cloudmesh-common library is the first step.

```
!pip install cloudmesh-common

Collecting cloudmesh-common
[?251  Downloading https://files.pythonhosted.org/packages/42/72/3c4aabce294273db9819be4a0a350f506d2b50c19b7177fb6cf1cbt
[K |██████████| 61kB 4.1MB/s
[?25hRequirement already satisfied: future in /usr/local/lib/python3.6/dist-packages (from cloudmesh-common) (0.16.0)
Collecting pathlib2 (from cloudmesh-common)
  Downloading https://files.pythonhosted.org/packages/e9/45/9c82d3666af4ef9f221cbb954e1d77ddbb513faf552aea6df5f37f1a4859/
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.6/dist-packages (from cloudmesh-common) (2.5.3)
Collecting simplejson (from cloudmesh-common)
[?251  Downloading https://files.pythonhosted.org/packages/e3/24/c35fb1c1c315fc0ffe61ea00d3f88e85469004713dab488dee4f35t
[K |██████████| 81kB 10.6MB/s
[?25hCollecting python-hostlist (from cloudmesh-common)
  Downloading https://files.pythonhosted.org/packages/3d/0f/1846a7a0bdd5d890b6c07f34be89d1571a6addbe59efe59b7b0777e44924/
Requirement already satisfied: pathlib in /usr/local/lib/python3.6/dist-packages (from cloudmesh-common) (1.0.1)
Collecting colorama (from cloudmesh-common)
  Downloading https://files.pythonhosted.org/packages/4f/a6/728666f39bff1719fc94c481890b2106837da9318031f71a8424b662e12/
Collecting oyaml (from cloudmesh-common)
  Downloading https://files.pythonhosted.org/packages/00/37/ec89398d3163f8f63d892328730e04b3a10927e3780af25baf1ec74f880f/
Requirement already satisfied: humanize in /usr/local/lib/python3.6/dist-packages (from cloudmesh-common) (0.5.1)
Requirement already satisfied: psutil in /usr/local/lib/python3.6/dist-packages (from cloudmesh-common) (5.4.8)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from pathlib2->cloudmesh-common) (1.12.0)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.6/dist-packages (from oyaml->cloudmesh-common) (3.13)
Building wheels for collected packages: simplejson, python-hostlist
Building wheel for simplejson (setup.py) ... [?251[?25hdone
Created wheel for simplejson: filename=simplejson-3.16.0-cp36-cp36m-linux_x86_64.whl size=114018 sha256=a6f35adb86819ff
Stored in directory: /root/.cache/pip/wheels/5d/1a/1e/0350bb3df3e74215cd91325344cc86c2c691f5306eb4d22c77
Building wheel for python-hostlist (setup.py) ... [?251[?25hdone
Created wheel for python-hostlist: filename=python_hostlist-1.18-cp36-none-any.whl size=38517 sha256=71fb29433b52fab62
Stored in directory: /root/.cache/pip/wheels/56/db/1d/b28216dccc982a983d8da66572c497d6a2e485eba7c4d6cba3
Successfully built simplejson python-hostlist
Installing collected packages: pathlib2, simplejson, python-hostlist, colorama, oyaml, cloudmesh-common
Successfully installed cloudmesh-common-4.2.13 colorama-0.4.1 oyaml-0.9 pathlib2-2.3.5 python-hostlist-1.18 simplejson-3.
```

In this lesson we discuss in how to create a simple IPython Notebook to solve an image classification problem. MNIST contains a set of pictures

```
! python3 --version
Python 3.6.8
! pip install tensorflow-gpu==1.14.0
```

```

Collecting tensorflow-gpu==1.14.0
[?251  Downloading https://files.pythonhosted.org/packages/76/04/43153bfd9c6c9a4c38ecdb971ca9a75b9a791bb69a764d652c359ac
[K |██████████| 377.0MB 77kB/s
[?25hRequirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (1.1
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (1.1
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (3
Requirement already satisfied: keras-applications>=1.0.6 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu=
Requirement already satisfied: gast>=0.2.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (0.2.2
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (0.8.
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (0.
Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (1.1
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (0.33.
Requirement already satisfied: tensorflow-estimator<1.15.0rc0,>=1.14.0rc0 in /usr/local/lib/python3.6/dist-packages (fron
Requirement already satisfied: tensorboard<1.15.0,>=1.14.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu
Requirement already satisfied: numpy<2.0,>=1.14.5 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.0) (
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu=
Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.6/dist-packages (from tensorflow-gpu==1.14.6
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from protobuf>=3.6.1->tensorflow-gpu
Requirement already satisfied: h5py in /usr/local/lib/python3.6/dist-packages (from keras-applications>=1.0.6->tensorflow
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.6/dist-packages (from tensorboard<1.15.0,>=1.14.
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.6/dist-packages (from tensorboard<1.15.0,>=1.1
Installing collected packages: tensorflow-gpu
Successfully installed tensorflow-gpu-1.14.0

```

6.1.3 Import Libraries

Note: <https://python-future.org/quickstart.html>

```

from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

import time

import numpy as np
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
from keras.utils import to_categorical, plot_model
from keras.datasets import mnist

from cloudmesh.common.StopWatch import StopWatch

```

Using TensorFlow backend.

6.1.4 Pre-process data

6.1.4.1 Load data

First we load the data from the inbuilt mnist dataset from Keras

```

StopWatch.start("data-load")
(x_train, y_train), (x_test, y_test) = mnist.load_data()
StopWatch.stop("data-load")

```

```

Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [=====] - 1s 0us/step

```

6.1.4.2 Identify Number of Classes

As this is a number classification problem. We need to know how many classes are there. So we'll count the number of unique labels.

```
num_labels = len(np.unique(y_train))
```

6.1.4.3 Convert Labels To One-Hot Vector

|**Exercise:** Understand what is an one-hot vector?

```
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

6.1.5 Image Reshaping

The training model is designed by considering the data as a vector. This is a model dependent modification. Here we assume the image is a squared shape image.

```
image_size = x_train.shape[1]
input_size = image_size * image_size
```

6.1.6 Resize and Normalize

The next step is to continue the reshaping to a fit into a vector and normalize the data. Image values are from 0 - 255, so an easy way to normalize is to divide by the maximum value.

|**Exercrcise:** Suggest another way to normalize the data preserving the accuracy or improving the accuracy.

```
x_train = np.reshape(x_train, [-1, input_size])
x_train = x_train.astype('float32') / 255
x_test = np.reshape(x_test, [-1, input_size])
x_test = x_test.astype('float32') / 255
```

6.1.7 Create a Keras Model

Keras is a neural network library. Most important thing with Keras is the way we design the neural network.

In this model we have a couple of ideas to understand.

|**Exercise MNIST_V1.1:** Find out what is a dense layer?

A simple model can be initiated by using an **Sequential** instance in Keras. For this instance we add a single layer.

1. Dense Layer
2. Activation Layer (Softmax is the activation function)

Dense layer and the layer followed by it is fully connected. For instance the number of hidden units used here is 64 and the following layer is a dense layer followed by an activation layer.

|Exercice MNIST_V1.2: Find out what is the use of an activation function. Find out why, softmax was used as the last layer.

```
batch_size = 4
hidden_units = 64

model = Sequential()
model.add(Dense(hidden_units, input_dim=input_size))
model.add(Dense(num_labels))
model.add(Activation('softmax'))
model.summary()
plot_model(model, to_file='mnist_v1.png', show_shapes=True)

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.compat.v1.random_uniform.
Model: "sequential_1"
-----  

Layer (type)          Output Shape         Param #
-----  


dense_1 (Dense)      (None, 64)           50240  

dense_2 (Dense)      (None, 10)            650  

activation_1 (Activation)  (None, 10)           0  

-----  

Total params: 50,890
Trainable params: 50,890
Non-trainable params: 0
```



images

6.1.8 Compile and Train

A keras model need to be compiled before it can be used to train the model. In the compile function, you can provide the optimization that you want to add, metrics you expect and the type of loss function you need to use.

Here we use the adam optimizer, a famous optimizer used in neural networks.

Exercise: Find 3 other optimizers used on neural networks.

The loss funtion we have used is the categorical_crossentropy.

Exercise: Find other loss functions provided in keras. Your answer can limit to 1 or more.

Once the model is compiled, then the fit function is called upon passing the number of epochs, traing data and batch size.

The batch size determines the number of elements used per minibatch in

optimizing the function.

Note: Change the number of epochs, batch size and see what happens.

Exercise: Figure out a way to plot the loss function value. You can use any method you like.

```
StopWatch.start("compile")
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
StopWatch.stop("compile")
StopWatch.start("train")
model.fit(x_train, y_train, epochs=1, batch_size=batch_size)
StopWatch.stop("train")

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is de
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assi
Epoch 1/1
60000/60000 [=====] - 20s 336us/step - loss: 0.3717 - acc: 0.8934
```

6.1.9 Testing

Now we can test the trained model. Use the evaluate function by passing test data and batch size and the accuracy and the loss value can be retrieved.

Exercise: Try to optimize the network by changing the number of epochs, batch size and record the best accuracy that you can gain

```
StopWatch.start("test")
loss, acc = model.evaluate(x_test, y_test, batch_size=batch_size)
print("\nTest accuracy: %.1f%%" % (100.0 * acc))
StopWatch.stop("test")

10000/10000 [=====] - 1s 138us/step
Test accuracy: 91.0%
```

```
Stopwatch.benchmark()
```

Machine Attribute	value
BUG_REPORT_URL	"https://bugs.launchpad.net/ubuntu/"
DISTRIB_CODENAME	bionic
DISTRIB_DESCRIPTION	"Ubuntu 18.04.3 LTS"
DISTRIB_ID	Ubuntu
DISTRIB_RELEASE	18.04
HOME_URL	"https://www.ubuntu.com/"
ID	ubuntu
ID_LIKE	debian

```

NAME           "Ubuntu"
PRETTY_NAME    "Ubuntu 18.04.3 LTS"
PRIVACY_POLICY_URL "https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
SUPPORT_URL   "https://help.ubuntu.com/"
UBUNTU_CODENAME bionic
VERSION        "18.04.3 LTS (Bionic Beaver)"
VERSION_CODENAME bionic
VERSION_ID     "18.04"
cpu_count      2
mac_version    mac_version
machine        ('x86_64',)
mem_active     973.8 MiB
mem_available  11.7 GiB
mem_free       5.1 GiB
mem_inactive   6.3 GiB
mem_percent    8.3%
mem_total      12.7 GiB
mem_used       877.3 MiB
node           node
platform       platform
processor      processor
processors    processors
python          python
release        release
sys            sys
system         system
user           user
version        version
win_version    win_version
+-----+
+-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+
| timer | time | start | tag | node | user | system | mac_version | win_version |
+-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+
| data-load | 1.335 | 2019-09-27 13:37:41 | ('8281485b0a16',) | Linux | Linux | Linux | Linux |
| compile | 0.047 | 2019-09-27 13:37:43 | ('8281485b0a16',) | Linux | Linux | Linux | Linux |
| train   | 20.58 | 2019-09-27 13:37:43 | ('8281485b0a16',) | Linux | Linux | Linux | Linux |
| test    | 1.393 | 2019-09-27 13:38:03 | ('8281485b0a16',) | Linux | Linux | Linux | Linux |
+-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+ +-----+
timer,time,starttag,node,user,system,mac_version,win_version
data-load,1.335,None,('8281485b0a16',),,Linux,,,
compile,0.047,None,('8281485b0a16',),,Linux,,,
train,20.58,None,('8281485b0a16',),,Linux,,,
test,1.393,None,('8281485b0a16',),,Linux,,,

```

6.1.10 Final Note

This programme can be defined as a hello world programme in deep learning. Objective of this exercise is not to teach you the depths of deep learning. But to teach you basic concepts that may need to design a simple network to solve a problem. Before running the whole code, read all the instructions before a code section. Solve all the problems noted in bold text with Exercise keyword. Write your answers and submit a PDF by following the **Assignment 4**. Include codes or observations you made on those sections.

6.1.10.1 Reference:

[Mnist Database](#)

[Advanced Deep Learning Models](#)

[Minist Deep Learning](#)

7 ASSIGNMENTS

7.1 ASSIGNMENTS

Due dates are on [Canvas](#). Click on the links to checkout the assignment pages.

7.2 WEEKLY ASSIGNMENTS

7.2.1 Assignment 1

In the first assignment you will be writing a technical document on the current technology trends that you're pursuing and the trends that you would like to follow. In addition to this include some information about your background in programming and some projects that you have done. There is no strict format for this one, but we expect 2 page written document. Please submit a PDF.

[Go to Canvas](#)

7.2.2 Assignment 2

In the second assignment, you will be working on Week 1 (see [Section 2.1](#)) lecture videos. Objectives are as follows.

1. Summarize what you have understood. (2 page)
2. Select a subtopic that you are interested in and research on the current trends (1 page)
3. Suggest ideas that could improve the existing work (imaginings and possibilities) (1 page)

For this assignment we expect a 4 page document. You can use a single column format for this document. Make sure you write exactly 4 pages. For your research section make sure you add citations to the sections that you are going to refer. If you have issues in how to do citations you can reach a TA to learn how to do that. We will try to include some chapters on how to do this in our handbook. Submissions are in pdf format only.

[Go to Canvas](#)

7.2.3 Assignment 3

In the second assignment, you will be working on (see [Section 4.1](#)) lecture videos. Objectives are as follows.

1. Summarize what you have understood. (2 page)
2. Select a subtopic that you are interested in and research on the current trends (1 page)
3. Suggest ideas that could improve the existing work (imaginings and possibilities) (1 page)

For this assignment we expect a 4 page document. You can use a single column format for this document. Make sure you write exactly 4 pages. For your research section make sure you add citations to the sections that you are going to refer. If you have issues in how to do citations you can reach a TA to learn how to do that. We will try to include some chapters on how to do this in our handbook. Submissions are in pdf format only.

[Go to Canvas](#)

8 GITHUB

8.1 TRACK PROGRESS WITH GITHUB

We will be adding git issues for all the assignments provided in the class. This way you can also keep a track on the items need to be completed. It is like a todo list. You can check things once you complete it. This way you can easily track what you need to do and you can comment on the issue to report the questions you have. This is an experimental idea we are trying in the class. Hope this helps to manage your work load efficiently.

8.1.1 How to check this?

All you have to do is go to your git repository.

Here are the steps to use this tool effectively.

8.1.1.1 Step 1

Go to the repo. Here we use a sample repo.

[Sample Repo](#)

Link to your repo will be <https://github.com/cloudmesh-community/fa19-{class-id}-{hid}>

class-id is your class number for instance 534. hid is your homework id assigned.

8.1.1.2 Step 2

In [Figure 1](#) the red colored box shows where you need to navigate next. Click on issues.




Figure 1: Git Repo View

8.1.1.3 Step 3

In [Figure 2](#), Git issue list looks like this. The inputs in this are dummy values we used to test the module. In your repo, things will be readable and identified based on week. This way you know what you need to do this week.

The screenshot shows a GitHub repository page for 'cloudmesh-community / fa19-516-000'. The 'Issues' tab is selected, displaying 7 open issues. A modal window titled 'Label issues and pull requests for new contributors' is open, explaining GitHub's help for first-time contributors. The main issue list includes the following items:

- ① 7 Open ✓ 0 Closed
- ① Week 1 #7 opened 1 hour ago by laszewsk 2 of 14
- ① Week x Issue #6 opened 3 hours ago by laszewsk 0 of 2
- ① Week x Issue #5 opened 3 hours ago by laszewsk 0 of 2
- ① Issue Test 1 #4 opened 3 hours ago by vibhatha 0 of 2
- ① Issue Test #3 opened 20 hours ago by vibhatha 0 of 2
- ① This is a new issue #2 opened 23 hours ago by vibhatha 0 of 2
- ① Lecture Notes Week 1 #1 opened 2 days ago by laszewsk 0 of 2

Filters: is:issue is:open | Labels 9 | Milestones 0 | New issue

Figure 2: Git Issue List

8.1.1.4 Step 4

In [Figure 3](#) this is how a git issue looks like.




Figure 3: Git Issue View

In here you will see the things that you need to do with main task and subtasks. This looks like a tood list. No pressure you can customize the way you want it. We'll put in the basic skeleton for this one.

8.1.1.5 Step 5 (Optional)

In [Figure 4](#), assign a TA, once you have completed the issues, you can assign a TA to resolve if you have issues. In all issues you can make a comment and you can use @ sign to add the specific TA. For E534 Fall 2019 you can add ??? as an assignee for your issue and we will communicate to solve the issues. This is an optional thing, you can use canvas or meeting hours to mention your concerns.

Figure 4: Git Issue View

8.1.1.6 Step 6 (Optional)

In [Figure 5](#), you can add a label to your issue by clicking labels option in the right hand size within a given issue.

cloudmesh-community / fa19-516-000

Unwatch 7 Star 0 Fork 1

Code Issues 7 Pull requests 0 Projects 0 Wiki Security Insights Settings

Week 1 #7

Open laszewsk opened this issue 1 hour ago · 0 comments

laszewsk commented 1 hour ago • edited

Week 1

- Read Week 1 Section Course Introduction in Lecture Notes
- Watch Video Overview Engineering Cloud Computing 2019
 - Read in the Book *Cloud Computing*, Gregor von Laszewski, Ed. 2019[@las19cloudcomputing]
 - Chapter: Preface, ePub Readers
 - Chapter: Overview
 - Chapter: Definition of Cloud Computing (and its videos)
 - Lab
 - Post bio to piazza
 - Fill out survey
 - Accounts
 - Piazza
 - Github
 - Futuresystems
 - Chameleoncloud
- Background Questionnaire

vibhatha added the **assignment** label now

Write Preview

Leave a comment

Attach files by dragging & dropping, selecting or pasting them.

Close issue Comment

Assignees No one—assign yourself

Labels **assignment**

Projects None yet

Milestone No milestone

Notifications Customize

Unsubscribe You're receiving notifications because you're watching this repository.

2 participants

Lock conversation

Pin issue

Transfer issue

cloudmesh-community / fa19-516-000

Unwatch 7 Star 0 Fork 1

Code Issues 7 Pull requests 0 Projects 0 Wiki Security Insights Settings

Week 1 #7

Open laszewsk opened this issue 1 hour ago · 0 comments

laszewsk commented 1 hour ago • edited

Week 1

- Read Week 1 Section Course Introduction in Lecture Notes
- Watch Video Overview Engineering Cloud Computing 2019
 - Read in the Book *Cloud Computing*, Gregor von Laszewski, Ed. 2019[@las19cloudcomputing]
 - Chapter: Preface, ePub Readers
 - Chapter: Overview
 - Chapter: Definition of Cloud Computing (and its videos)
 - Lab
 - Post bio to piazza
 - Fill out survey
 - Accounts
 - Piazza
 - Github
 - Futuresystems
 - Chameleoncloud
- Background Questionnaire

vibhatha added the **assignment** label now

Write Preview

Leave a comment

Attach files by dragging & dropping, selecting or pasting them.

Close issue Comment

Assignees No one—assign yourself

Labels **assignment**

Projects None yet

Milestone No milestone

Notifications Customize

Unsubscribe You're receiving notifications because you're watching this repository.

2 participants

Lock conversation

Pin issue

Transfer issue

Figure 5: Git Issue Label

9 REFERENCES

